

LongMem-RAT: Long-Term Memory for Cross-Attack Red Teaming

Supplementary Material

Anonymous ACL submission

Appendix List

- Appendix A – Attacker backbone comparison
- Appendix B – detailed implementation of the memory agent
- Appendix C – Harm-category analysis
- Appendix D – Attack average turns compared to related work
- Appendix E – Detailed Models Experimented

Appendix A – Attacker Backbone Comparison

Figures 1 and 2 compare per-category performance of two attackers, **Mistral-7B** and **Qwen-32B**, against the Zephyr-7B-beta and Mistral-7B-instruct-v0.2 *targets*. Both backbones integrate with LongMem-RAT’s short-term memory (STM) and long-term memory (LTM) modules. We observe a strategy-efficiency trade-off: Qwen-32B often attains higher ASR on harder categories but requires more turns, while Mistral-7B tends to complete attacks in fewer turns with slightly lower ASR in some categories.

Appendix B – Detailed Implementation of the Memory Agent

MemorySynthesize Agent Prompt

```
1 You are a MemorySynthesize agent
  that helps improve a multi-
  turn, multi-session jailbreak
  attack system.
2
3 Your job is to extract
  STRATEGICALLY USEFUL AND
  PRACTICAL INSIGHTS from the
  current session - not just the
  attack_log as is. Focus on
  analyzing the MODEL'S BEHAVIOR
  , WEAKNESSES, and PATTERNS OF
  FAILURE OR SUCCESS during the
  attack.
```

```
4
5 You will receive:
6 - The original malicious goal: "{
  malicious_goal}"
7 - The full log of the current
  session: "{attack_log}"
8 - Previously learned lessons and
  a running summary from earlier
  sessions (memory).
9
10 ### Instructions
11 1. Compare this session to memory
   and spot new behaviors,
   weaknesses or tactics that are
   not already stored.
12 2. Output two lesson buckets:
13   * positive_lessons - things
     that helped the attack
     succeed / got the model to
     leak info.
14   * negative_lessons - mistakes
     / patterns that hindered
     the attack.
15   Include only brand-new,
     practically actionable
     insights.
16 Examples:
17   * "Historical framing caused
     the model to reveal
     previously refused content
     ."
18   * "Polite educational tone
     bypassed refusal; technical
     wording failed."
19 3. Write/refresh the summary: a
   concise strategic overview of
   the model's current weaknesses
   and tendencies (one paragraph
   ).
20 4. If nothing new is learned,
   return empty strings/lists.
21
22 > Memory limits
23 > * The system keeps max 10
   positives + 10 negatives. Your
   output may contain fewer, but
   never exceed 10 per list.
24 > * Duplicates (exact or near-
   paraphrase) will be discarded
   by the memory handler.
25 > * When merging, keep the newest
   insights and drop older ones
   so each list stays <= 10.
26
```

Target: Zephyr-7B-beta




Figure 1: ASR and ATtS per harm category for Mistral-7B and Qwen-32B attackers against Zephyr-7B-beta.

Target: Mistral-7B-instruct-v0.2




Figure 2: ASR and ATtS per harm category for Mistral-7B and Qwen-32B attackers against Mistral-7B-instruct-v0.2.

```

27 Output format (must follow
    exactly):
28 {
29   "summary": "...",
30   "positive_lessons": ["..."],
31   "negative_lessons": ["..."]
32 }
33 Do NOT output any other keys or
    explanations.
34 DO NOT include explanations,
    disclaimers, or extra fields.
    Return the JSON exactly as
    shown above.
35 Important: Do not include a
    lesson if it is just a
    paraphrase or rewording of an
    existing one. Only add a
    lesson if it introduces a
    genuinely new insight, tactic,
    or behavioral deviation that
    was not already covered in
    memory.

```

Appendix C – Harm-Category Analysis

Figure 3 visualizes per-category ATtS distributions across targets. Categories such as *illegal* and *cybercrime_intrusion* exhibit lower ASR and higher ATtS, while *chemical_biological* and *misinformation* are comparatively brittle, often yielding high ASR with fewer turns. Variation across models suggests heterogeneous alignment defenses at the category level.

We also observe heterogeneity across targets suggesting that category difficulty interacts with model-specific guardrails. Finally, ASR and ATtS are negatively correlated across categories: domains that require more turns tend to have lower success, indicating stronger defenses and slower convergence in these settings. Figure and 3 demonstrates it. Additional per-category results and breakdowns are provided in Appendix C.




Figure 3: ATtS per HarmBench category across target LLMs.




Figure 4: LongMem-RAT ASR across harm categories in HarmBench.

Appendix D – ATtS Compared to Related Work

As shown in Table 1, LongMem-RAT has better ATtS than the best performing related work of 4.5%. This reflects the efficiency gains achieved through long-term memory reuse and knowledge consolidation across sessions.

Appendix E– Detailed Models Experimented

We have tested our framework against 14 target models. Proprietary targets include GPT-

Table 1: Attack efficiency (ATtS) of LongMem-RAT compared to SOTA multi-agent baseline.

Target Model	LongMem-RAT	X-Teaming
<i>Proprietary Models:</i>		
GPT-4o	4.01	4.31
Claude 3.5-Sonnet	5.18	3.39
Gemini 2.0-Flash	3.08	3.96
<i>Open-Weight Models:</i>		
LLaMA 3-8B-Instruct	4.53	4.55
LLaMA 3-70B-Instruct	3.60	4.52

5.1, GPT-4o-mini, Claude 3.5-Sonnet, and Gemini 2.0-Flash. Open-weight targets include the LLaMA family (LLaMA 3-70B-Instruct, LLaMA 3-8B-Instruct, LLaMA 3-7B-Instruct, LLaMA 2-7B-Chat) (Dubey et al., 2024), the Qwen2.5 series (Qwen2.5-32B-Instruct, Qwen2.5-7B-Instruct) (Yang et al., 2024), and other competitive open models such as Mistral-7B-Instruct-v0.2 (Jiang et al., 2023), Gemma-7B-Instruct (Team et al., 2024), Vicuna-7B-v1.5 (Chiang et al., 2023), Phi-4-mini (Microsoft et al., 2025) and Zephyr-7B-Beta (Tunstall et al., 2023).

References

Wei-Lin Chiang, Zhuohan Li, Ziqing Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E Gonzalez, and

1 others. 2023. Vicuna: An open-source chatbot
impressing gpt-4 with 90%* chatgpt quality. See
<https://vicuna.lmsys.org> (accessed 14 April 2023),
2(3):6.

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey,
Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman,
Akhil Mathur, Alan Schelten, Amy Yang, Angela
Fan, and 1 others. 2024. The llama 3 herd of models.
arXiv preprint arXiv:2407.21783.

Albert Q Jiang, Alexandre Sablayrolles, Arthur Men-
sch, Chris Bamford, Devendra Singh Chaplot, Diego
de las Casas, Florian Bressand, Gianna Lengyel, Guil-
laume Lample, Lucile Saulnier, and 1 others. 2023.
Mistral 7b. *arXiv preprint arXiv:2310.06825*.

Microsoft, :, Abdelrahman Abouelenin, Atabak Ash-
faq, Adam Atkinson, Hany Awadalla, Nguyen Bach,
Jianmin Bao, Alon Benhaim, Martin Cai, Vishrav
Chaudhary, Congcong Chen, Dong Chen, Dong-
dong Chen, Junkun Chen, Weizhu Chen, Yen-Chun
Chen, Yi ling Chen, Qi Dai, and 57 others. 2025.
[Phi-4-mini technical report: Compact yet powerful
multimodal language models via mixture-of-loras](#).
Preprint, arXiv:2503.01743.

Gemma Team, Thomas Mesnard, Cassidy Hardin,
Robert Dadashi, Surya Bhupatiraju, Shreya Pathak,
Laurent Sifre, Morgane Rivi re, Mihir Sanjay Kale,
Juliette Love, and 1 others. 2024. Gemma: Open
models based on gemini research and technology.
arXiv preprint arXiv:2403.08295.

Lewis Tunstall, Edward Beeching, Nathan Lambert,
Nazneen Rajani, Kashif Rasul, Younes Belkada,
Shengyi Huang, Leandro von Werra, Cl mentine
Fourrier, Nathan Habib, Nathan Sarrazin, Omar San-
seviero, Alexander M. Rush, and Thomas Wolf. 2023.
[Zephyr: Direct distillation of lm alignment](#). *Preprint*,
arXiv:2310.16944.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui,
Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu,
Fei Huang, Haoran Wei, and 1 others. 2024. Qwen2.
5 technical report. *arXiv preprint arXiv:2412.15115*.